

# CS7IS2 - Artificial Intelligence Assignment 1 Report

Kaaviya Paranji Ramkumar

March 2025

## 1 Introduction

This assignment presents the implementation and comparative analysis of various search and Markov Decision Process (MDP) algorithms for maze-solving problems. A maze generator capable of creating environments of different dimensions was developed to serve as the testing environment for the algorithm implementations. The search algorithms implemented include Depth-First Search (DFS), Breadth-First Search (BFS), and A\* to find paths through the mazes. Additionally, two MDP-based approaches - Value Iteration and Policy Iteration - were implemented to solve the same maze problems.

The core of this work involved the careful design of each algorithm to efficiently navigate through maze environments, followed by extensive testing across mazes of varying complexity. Performance data was collected using metrics such as execution time, memory usage, and solution quality to provide a comprehensive comparison between these different approaches. The analysis explores the strengths and limitations of each algorithm, examining not only how search algorithms compare to each other, but also how MDP-based approaches differ from traditional search methods in maze-solving tasks.

## 2 Methodology

### 2.1 Maze Generation

#### Description of the maze generator

The maze generation implementation utilizes a Recursive Backtracker algorithm with enforced boundaries. This approach generates perfect mazes that contain exactly one path between any two points, with no loops or inaccessible areas. The implementation creates a grid where cells are either walls ('#') or paths ('.'), with special markers for start ('S') and goal ('G') positions. Below are listed the key steps involved in implementing the maze generator:

- Initialize maze with custom size with all cells as walls.
- Create a NumPy grid with odd dimensions to ensure proper wall placement.
- Implement depth-first search with backtracking using a stack for finding path.
- Perform random selection of unvisited neighboring cells during path creation.
- Optionally, add multiple paths by randomly removing walls based on a configurable path density parameter.
- Enforce boundary walls that cannot be removed, so that the maze remains contained.

#### Maze properties

The implemented maze generator produces mazes with the following properties:

- Fully connected structure with enforced boundary walls.
- Default start position at (1,1) and goal position at the bottom-right corner (height-2, width-2).

- Configurable dimensions that are always adjusted to odd numbers to maintain proper wall-path structure.
- Optional path density parameter that can create shortcuts by removing interior walls. This in turn created multiple paths from start to goal position.

For testing purposes, a range of odd-sized mazes was generated:

- $7 \times 7$  (very small)
- $11 \times 11$ ,  $15 \times 15$ ,  $19 \times 19$ ,  $21 \times 21$ ,  $25 \times 25$  (small to medium)
- $31 \times 31$ ,  $45 \times 45$ ,  $59 \times 59$ ,  $67 \times 67$  (medium to large)
- $83 \times 83$ ,  $101 \times 101$  (large)

The maze complexity increases with size, as larger dimensions create longer potential solution paths and a larger state space for algorithms to explore.

#### Justification for chosen maze sizes

The selection of maze sizes for performance comparison was based on several key considerations:

- **Consistent structure:** All sizes are odd numbers, which aligns with the recursive backtracker algorithm’s requirement for proper wall placement between path cells.
- **Graduated complexity range:** The sizes range from very small ( $7 \times 7$ ) to quite large ( $101 \times 101$ ), providing a comprehensive spectrum to observe how algorithm performance scales with increasing problem complexity.
- **Finer granularity at critical transition points:** More sizes were selected in the small-to-medium range ( $11 \times 11$  through  $25 \times 25$ ) to capture the performance characteristics to see where algorithms begin to show meaningful differences in behavior.
- **Performance constraints consideration:** The largest sizes ( $83 \times 83$  and  $101 \times 101$ ) were included to test the limits of the algorithms and the maze generator itself, revealing important distinctions in time and memory efficiency under high computational demands.
- **Statistical significance:** Multiple maze instances were generated at each size to ensure results were not biased by any particular maze configuration, providing more reliable performance metrics.

This carefully selected range of maze sizes enables a thorough analysis of how different search and MDP algorithms respond to increasing problem complexity in terms of execution time, memory usage, and solution quality.

## 2.2 Algorithm Implementations

### 2.2.1 Search Algorithms

#### Depth-First Search (DFS) and Breadth-First Search (BFS)

Both the DFS and BFS implementations share similar design patterns, with key differences in their traversal strategies. These classic graph search algorithms were implemented with careful consideration of several important factors.

- **State representation:** For the representation of states, coordinates (x, y) within the maze grid were chosen as the primary way to track positions. This approach gives an intuitive mapping between the algorithm’s internal state and the physical maze structure. Using simple coordinate pairs makes it easy to calculate adjacent positions and requires minimal memory overhead, which becomes especially important when dealing with larger maze sizes.

- **Action space:** The action space was deliberately limited to the four cardinal movements: up, down, left, and right. This design prevents diagonal movements that might allow "cutting corners" through walls while ensuring complete coverage of all possible paths through the maze. Each movement has a uniform cost of 1 unit, simplifying path cost calculations and maintaining consistency across different maze configurations.
- **Goal recognition:** Goal recognition is handled through explicit marking with the 'G' character in the maze array. This approach makes it straightforward to detect when the search has reached the target position, supporting flexible goal placement anywhere within the maze structure. Such clear termination condition helps prevent unnecessary exploration once the goal has been located.
- **Fringe management:** The key difference between DFS and BFS lies in their fringe management strategies. DFS employs a stack data structure that encourages exploring as deeply as possible along each branch before backtracking. This can be advantageous in mazes with long corridors and few branching paths. In contrast, BFS utilizes a queue that ensures all states at a given depth are fully explored before moving deeper, guaranteeing the shortest path will be found when a solution exists.
- **Track visited states:** To prevent cycles and repeated exploration, both algorithms maintain a boolean matrix matching the maze dimensions to track visited states. This approach offers constant-time lookup performance and consumes less memory than alternative approaches like hash sets, especially for larger mazes where memory efficiency becomes critical.
- **Parent pointer map:** For path reconstruction, a parent pointer map is maintained throughout the search process. This dictionary maps each visited state to the state that led to its discovery. When the goal is reached, the solution path can be efficiently reconstructed by working backward through these parent pointers until reaching the starting position. This eliminates the need to store complete paths during the search phase, significantly reducing memory requirements.

## A\* Search

The A\* search algorithm builds upon the foundation of BFS but incorporates a heuristic function to guide the search toward the goal more efficiently. This approach combines the completeness of BFS with additional intelligence about the problem structure.

- **Heuristic function:** The Manhattan distance was selected as the heuristic function for our A\* implementation after careful consideration of alternatives. This heuristic is particularly well-suited for grid-based mazes with cardinal movements because it exactly represents the minimum number of steps required to reach the goal in the absence of walls. This property ensures the heuristic never overestimates the true cost, making it admissible and guaranteeing optimal solutions. Additionally, Manhattan distance satisfies the consistency requirement for A\*, meaning the estimated cost to the goal never increases by more than the actual cost of moving between states. From a computational standpoint, the calculation involves simple arithmetic operations, adding minimal overhead to each state evaluation.
- **Priority queue implementation:** To manage the exploration frontier efficiently, a priority queue was implemented using a three-part priority tuple: (f\_score, counter, position). The f\_score combines the known cost from the start (g\_score) with the estimated cost to the goal (heuristic), directing the search toward promising paths. The counter serves as a tiebreaker for states with equal f\_scores, ensuring consistent behavior and preventing potential oscillations between equally valued states. This approach avoids the need for costly priority queue updates when better paths to already-discovered states are found.
- **Path cost tracking:** Path costs are tracked using a dictionary that maps each explored state to its current best-known distance from the start. When the algorithm discovers a potentially better path to a state, it compares the new path cost with the previously recorded value and updates accordingly. This mechanism ensures that the search always maintains optimal path information throughout the exploration process, even when encountering the same state through different routes.

The implementation also includes comprehensive performance metrics collection to facilitate meaningful algorithm comparisons. These metrics include the number of nodes explored (providing insight into search efficiency), maximum fringe size (indicating peak memory usage), path length (measuring solution quality), and execution time (quantifying computational performance). These measurements allow for data-driven evaluation of how A\* performs relative to the uninformed search algorithms across different maze configurations and sizes.

## 2.2.2 Markov Decision Process (MDP) Algorithms

### Value Iteration and Policy Iteration

The Markov Decision Process (MDP) offers a different approach to maze solving compared to traditional search algorithms. Instead of explicitly searching for a path, MDP algorithms compute an optimal policy that specifies the best action to take at each state. Two classic MDP algorithms were implemented: Value Iteration and Policy Iteration.

For the MDP formulation of the maze-solving problem, these are the common key design decisions were made:

- **State representation:** States are represented as coordinates within the maze, similar to the search algorithms. This representation naturally maps to the grid structure while keeping the state space manageable and allowing for direct performance comparisons.
- **Action space:** The action space consists of the four cardinal directions (up, down, left, right), maintaining consistency with the search implementations and realistically modeling movement constraints in a grid-based maze.
- **Rewards:** The reward structure was carefully designed to guide the agent toward the goal while discouraging excessive movements. A large positive reward (+100) is assigned to reaching the goal state, creating a strong incentive to find it. Small negative rewards are assigned to all intermediate states, effectively creating a path cost that encourages finding shorter paths. Additionally, attempting to move into a wall results in remaining in the same state and still incurring the negative reward (-100), discouraging such attempts.
- **Transition model:** The transition model is deterministic in our implementation, meaning that taking an action from a state always leads to the same next state with probability 1.0. This design choice simplifies the MDP while still capturing the essential dynamics of maze navigation and creates a fair comparison with the deterministic search algorithms.
- **Discount factor:** A discount factor  $\gamma$  (gamma) between 0.9 and 0.99 was selected after experimentation with different values. This parameter balances the importance of immediate versus future rewards. Setting gamma close to 1 encourages the agent to consider long-term rewards more heavily, which is appropriate for maze-solving where the goal might be many steps away.

For the Value Iteration algorithm, individual key implementation details include:

- An adaptive convergence threshold that ensures the algorithm converges to an approximately optimal policy while avoiding unnecessary computation.
- Iterative computation of state values until the maximum change between iterations falls below the threshold.
- A threshold of 0.001 that provides a good balance between solution quality and computational efficiency

As for the Policy Iteration, algorithm alternates between policy evaluation and policy improvement steps. For policy evaluation, a system of linear equations is solved to determine the values of states under the current policy. Policy improvement then updates the policy based on these values. This two-step process often converges in fewer iterations than Value Iteration, especially for larger mazes, though each iteration typically requires more computation.

Both MDP algorithms were implemented with a custom maximum iteration limit to ensure termination even for very large or complex mazes. Based on experimentation with different maze sizes, a limit of 1000 iterations was found to be more than sufficient for convergence in tested mazes.

To extract a solution path once the optimal policy has been computed, a path extraction method follows the optimal policy from the initial state until reaching the goal state. This allows direct comparison of path lengths with the search algorithms and provides a concrete solution rather than just a policy.

Performance metrics calculated for both MDP algorithms include:

- Convergence iterations
- Execution time
- Solution path length
- Change in state values between iterations

## 2.3 Experimental Setup

### Code organization and object-oriented design

The experimental framework for this assignment was carefully planned and implemented using object-oriented programming principles to ensure modularity, re-usability, and clean separation of concerns. The code structure follows software engineering best practices, making it easy to extend and maintain.

The code base for maze generation and solving is currently maintained at <https://github.com/prkaaviya/TCD-Artificial-Intelligence-A1>

The project is organized in a hierarchical structure with clearly defined components like:

```
.
|-- Makefile           # For automation
|-- gen_maze.py        # Maze generation module
|-- run.py             # Main execution script to solve maze
|-- mazes/            # Generated maze data
|   |-- text/         # Text representation of mazes
|   |-- visuals/      # Visual representation of mazes
|-- notebooks/        # Analysis and debugging notebooks
|-- results/          # Algorithm output
|   |-- metrics/      # Performance measurements
|   |-- visuals/      # Solution visualizations
|-- solvers/          # Algorithm implementations
|   |-- base.py       # Base solver class
|   |-- informed.py   # A* implementation
|   |-- mdp.py        # MDP algorithm implementations
|   |-- uninformed.py # DFS and BFS implementations
|   |-- utils.py      # Utility functions
```

- **Solver classes:** All maze-solving algorithms are implemented as classes that inherit from a common base class (`MazeSolverBase` in `solvers/base.py`). This base class provides shared functionality such as maze loading, visualization, and performance metrics tracking. Using inheritance ensures consistent interfaces across all algorithm implementations while allowing each algorithm to have its specialized behavior.
- **Algorithm implementations:** The algorithms are logically grouped into separate modules:
  - `uninformed.py` contains DFS and BFS implementations.
  - `informed.py` contains the A\* implementation.
  - `mdp.py` contains Value Iteration and Policy Iteration implementations.

## Automation with Makefile

A comprehensive Makefile was developed to automate all aspects of the experimentation process, providing a user-friendly interface for running multiple tests and saving overall time taken to conduct testing and analysis. The Makefile includes targets for:

- Maze generation commands to generate mazes of various sizes ( $7 \times 7$  to  $101 \times 101$ ) with consistent properties.
- Commands to run specific algorithms on specific maze sizes.
- Execute benchmarking with commands to automatically run all algorithms on all maze sizes.

Such automation proved crucial for ensuring consistent testing conditions across all algorithms and maze sizes. The Makefile also manages the organization of output files, storing results in a structured directory hierarchy that separates metrics data and visualizations.

For example, running `make benchmark` automatically executes all five algorithms (DFS, BFS, A\*, Value Iteration, and Policy Iteration) on all available maze sizes, generating comprehensive performance data in a single command. This approach minimizes human error in the experimental process and ensures reproducibility of results.

### Maze generation outputs

The maze generation process produces two distinct outputs for each maze size:

- A text file representation of the maze stored as a NumPy array in CSV format (.txt), where characters represent different maze elements: '#' for walls, '.' for open paths, 'S' for start position, and 'G' for goal position. This format allows for easy loading and processing by the solver algorithms.
- A visual representation of the maze saved as a PNG image, where black cells represent walls, white cells represent paths, green marks the start position, and red indicates the goal. These visualizations provide an intuitive way to understand the maze structure and complexity.

All generated mazes are stored in a structured directory hierarchy, with text representations in `mazes/text/` and visual representations in `mazes/visuals/`. This organization ensures that mazes can be consistently referenced by both the solver algorithms and the analysis tools.

### Data Collection and Metrics

The framework systematically collects a comprehensive set of performance metrics for each algorithm:

- *execution\_time* is measured in seconds, capturing the computational efficiency of each algorithm.
- *nodes\_explored* tracks how many states each algorithm visits during execution.
- *path\_length* measures the quality of the solution found.
- *iterations* metric is available for MDP algorithms, which tracks how many iterations were needed for convergence.

These metrics are automatically saved to CSV files in the `results/metrics` directory, with naming conventions that clearly identify the maze size and algorithm used. Additionally, visual representations of the solution paths are rendered and saved to the `results/visuals` directory, providing intuitive visualization of how each algorithm navigates the maze.

Such meticulous experimental setup provides a robust foundation for analyzing the performance characteristics of different search and MDP algorithms, ensuring that conclusions drawn from the results are based on systematic and reproducible testing.

## 3 Results and Analysis

### 3.1 Search Algorithm Comparison (DFS, BFS, A\*)

When comparing DFS, BFS, and A\* algorithms against each other, several key performance differences emerge based on the collected metrics that are broadly discussed below.

- **Path quality:** As we can see from “Fig. 1“, BFS guarantees the shortest path in unweighted graphs, while DFS often produces inefficient paths as it explores deeply before backtracking. At the largest maze size (101), DFS paths are approximately 8 times longer than optimal. A\* consistently finds optimal paths identical to BFS but does so more efficiently by using heuristic guidance.
- **Computational efficiency:** Execution time analysis from “Fig. 2“ reveals that all three algorithms have similar asymptotic growth as maze sizes increase, with execution times ranging from 0.1 to 10ms. A\* shows interesting performance fluctuations at sizes 21 and 67, we can infer its heuristic might perform differently depending on maze structure. Despite these fluctuations, all algorithms seem to maintain comparable execution times across maze sizes, with slight advantages varying by maze complexity.
- **Memory usage:** “Fig. 3“ demonstrates the peak memory usage during maze solving for each algorithm and shows that DFS consumes significantly more memory than both BFS and A\*, particularly in larger mazes. At size 101, DFS requires approximately 4.4MB, while BFS uses only about 0.15MB and A\* about 0.6MB. This actually contradicts the theoretical expectation that BFS would use more memory, suggesting implementation details might be affecting the measurements or that the maze structure with multiple paths from start to goal can influence memory patterns in the case of DFS.
- **Nodes explored:** BFS explores the most nodes across almost all maze sizes, followed by DFS, with A\* exploring the fewest as seen in “Fig. 4“. At size 101, BFS examines approximately 5400 nodes, DFS around 4400, and A\* only about 4000. This demonstrates A\*'s efficiency in avoiding unnecessary exploration. The data also shows an interesting anomaly at size 67 where A\* explores significantly fewer nodes, suggesting maze structure may particularly favor A\*'s heuristic in certain configurations.
- **Scalability:** As maze dimensions increase, A\* demonstrates superior scalability by maintaining the optimal path length while exploring fewer nodes than both DFS and BFS. The logarithmic growth in execution time across all algorithms suggests good scalability, though with different trade-offs: DFS sacrifices path quality for potentially better space efficiency in some implementations, BFS guarantees optimality but explores more nodes, and A\* achieves optimality while minimizing exploration.

**Why does BFS seem to be more memory efficient than DFS?** The theoretical space complexity of DFS is  $O(d)$  and BFS is  $O(bd)$ , where  $d$  is the depth of the solution and  $b$  is the branching factor. However, the current implementation shows that practical memory usage can diverge significantly from theoretical predictions due to specific implementation decisions.

The primary factor driving memory consumption in my implementation is the decision to store complete path histories with each frontier node. While this approach simplifies solution path reconstruction, it fundamentally alters the memory usage patterns:

- In DFS, the paths stored with each node grow rapidly as the algorithm explores deeply, creating a multiplicative effect on memory usage. Since DFS paths are significantly longer especially in large maze size (approximately 8 times longer in  $101 \times 101$  mazes), the memory overhead grows accordingly.
- In BFS, while more total nodes may be explored, the paths associated with each node remain relatively short until the goal is reached. This results in substantially lower memory requirements despite the theoretically higher space complexity.

Additionally, maze complexity can also significantly impact memory consumption in the case of DFS:

- As maze size and complexity increase, DFS tends to explore deeper into “dead-end” paths before backtracking, leading to excessive path lengths as clearly visible in “Fig. 5” where BFS and DFS solution paths are visualized on the maze.




Figure 1: Comparison of path lengths across different maze sizes for search algorithms.

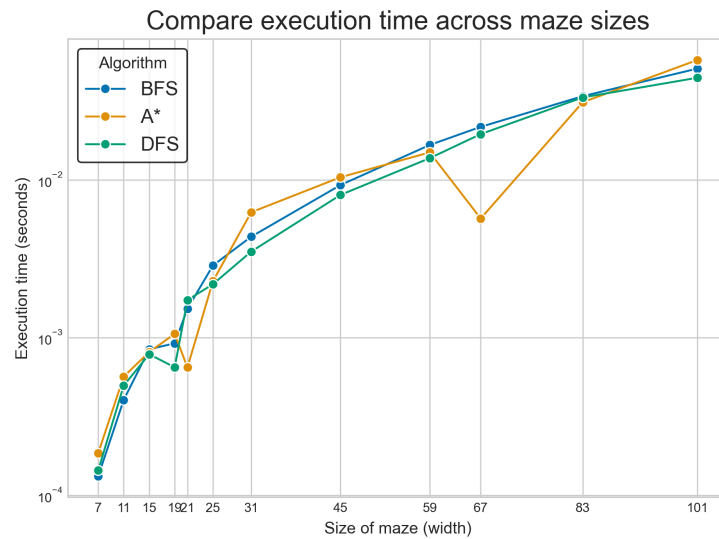


Figure 2: Comparison of execution time across different maze sizes for search algorithms.






Figure 3: Comparison of peak memory usage across different maze sizes for search algorithms.





Figure 4: Plot of number of nodes explored for different maze sizes with search algorithms.

maze45 (45, 45) Solution with BFS




(a) Maze  $45 \times 45$ , BFS solution

maze45 (45, 45) Solution with DFS




(b) Maze  $45 \times 45$ , DFS solution

maze59 (59, 59) Solution with BFS




(c) Maze  $59 \times 59$ , BFS solution

maze59 (59, 59) Solution with DFS




(d) Maze  $59 \times 59$ , DFS solution

maze67 (67, 67) Solution with BFS



(e) Maze  $67 \times 67$ , BFS solution


maze67 (67, 67) Solution with DFS



(f) Maze  $67 \times 67$ , DFS solution


Figure 5: Comparison of BFS and DFS solutions for mazes of sizes  $45 \times 45$ ,  $59 \times 59$ , and  $67 \times 67$ . The left column shows BFS solutions, which find optimal (shortest) paths. The right column shows DFS solutions, which tend to explore more of the maze and produce longer paths.

maze83 (83, 83) Solution with BFS




(g) Maze  $83 \times 83$ , BFS solution

maze83 (83, 83) Solution with DFS




(h) Maze  $83 \times 83$ , DFS solution

maze101 (101, 101) Solution with BFS



(i) Maze  $101 \times 101$ , BFS solution

maze101 (101, 101) Solution with DFS



(j) Maze  $101 \times 101$ , DFS solution

Figure 5: Comparison of BFS and DFS solutions for mazes of sizes  $83 \times 83$  and  $101 \times 101$ . The difference in path length becomes increasingly pronounced as maze size increases.

- The side-by-side comparison of solution paths (particularly in the  $83 \times 83$  and  $101 \times 101$  mazes) visually demonstrates how DFS paths meander through large portions of the maze, while BFS paths take direct routes.
- Each additional turn or segment in these winding DFS paths compounds the memory overhead, as every node added to the stack carries a copy of an increasingly long path history.
- In complex mazes with many branching paths, this path-copying behavior creates a multiplicative effect on memory usage that grows more severe with maze size, explaining the steep upward trend in DFS memory consumption shown in "Fig. 3".

This finding highlights an important principle in algorithm implementation: theoretical space complexity assumes minimal state representation, while practical implementations often include additional state information that can significantly alter performance characteristics. When path reconstruction is a requirement, the traditional space advantage of DFS might get negated by the longer paths it generates.

### 3.2 MDP Algorithm Comparison (Value Iteration vs Policy Iteration)

The performance comparison between Value Iteration and Policy Iteration algorithms for solving Markov Decision Process models of the maze reveals several key differences and huge challenges specific to MDP-based approaches.

- **Solution convergence:** Both algorithms successfully found solutions for smaller mazes, converging to optimal policies. Value Iteration typically required more iterations to converge than Policy Iteration, as evidenced in their respective iteration counts plotted in "Fig. 6". For instance, in this maze solver, Value Iteration often needed 1.5-2x more iterations than Policy Iteration to reach convergence thresholds.
- **Computational efficiency:** A critical finding was the significant computational burden of MDP algorithms on larger mazes as indicated by the time taken to solve the maze in "Fig. 7". While search algorithms (DFS, BFS, A\*) completed even for  $101 \times 101$  mazes in milliseconds, MDP algorithms became prohibitively expensive beyond certain maze dimensions. Policy Iteration particularly suffered from the cubic complexity of its policy evaluation step, making it computationally infeasible for mazes larger than approximately  $45 \times 45$  cells in my testing environment even after reducing the discount factor.
- **Memory usage:** Both MDP algorithms showed similar memory consumption patterns, proportional to the number of states in the maze as depicted in "Fig. 8". However, their memory requirements were substantially higher than search algorithms on an average for equivalent maze sizes due to the need to maintain complete value functions and policies across all states simultaneously.
- **State evaluations:** "Fig. 9" reveals exponential growth in state evaluations as maze sizes increase, with Policy Iteration consistently requiring 5-10 times more evaluations than Value Iteration. Both algorithms fail to find solutions at larger maze sizes ( $45 \times 59$ ), despite evaluating up to  $10^6$  states. This shows a fundamental scalability issue with MDP approaches in maze-solving, especially as performance degrades significantly beyond  $31 \times 31$  mazes.
- **Parameter sensitivity:** Both algorithms showed significant sensitivity to parameter choices. The discount factor ( $\gamma$ ) particularly impacted convergence speed and solution quality. With higher values (e.g., 0.999) creating precise but computationally expensive solutions and sometimes no solution at all, and lower values (e.g., 0.8) offering faster convergence at the cost of potential solution quality.

#### MDP Limitation with scalability

A fundamental observation from our experiments is the severe scalability limitation of MDP approaches compared to search algorithms:




Figure 6: Comparison of iterations required for convergence across different maze sizes for MDP algorithms.

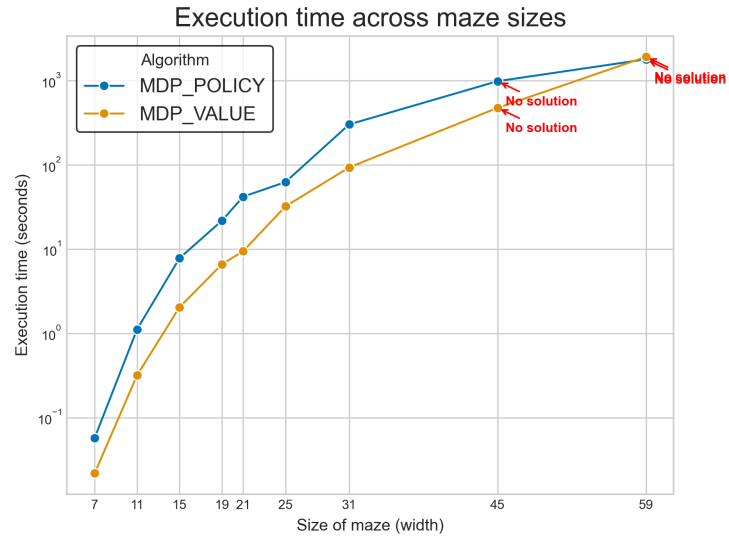


Figure 7: Comparison of execution time across different maze sizes for MDP algorithms, showing the dramatic increase in computational cost as maze sizes increase.




Figure 8: Comparison of memory consumed across different maze sizes for MDP algorithms.




Figure 9: Comparison of states evaluated across different maze sizes for MDP algorithms, showing the huge increase in evaluation stage as maze sizes increase.




Figure 10: Plot of amount of policy changes executed during maze solving with MDP Policy algorithm.

- **State space explosion:** MDP algorithms suffer from the "curse of dimensionality" where computational requirements grow exponentially with the number of states. For this maze problem, it translated to quadratic growth in states with respect to maze dimensions (width $\times$ height).
- **Computation time growth:** While search algorithms showed moderate growth in execution time (from 0.1ms to 10ms) as maze size increased from 7 $\times$ 7 to 101 $\times$ 101, MDP algorithms exhibited dramatic increases from milliseconds for small mazes to several minutes or even hours for mazes beyond 45 $\times$ 45. Several larger maze configurations were unable to find solution and/or complete within practical time constraints.
- **Value Iteration vs Policy Iteration:** Value Iteration performed marginally better scalability than Policy Iteration for larger mazes, likely due to Policy Iteration's need to solve systems of equations during policy evaluation. However, both algorithms implementation became impractical beyond certain maze dimensions.

This analysis reveals that while MDP methods provide a theoretically sound approach to maze-solving with guaranteed optimality, their practical application is limited to smaller maze instances due to computational constraints. For large-scale maze solving, search algorithms like A\* offer a significantly more efficient alternative while still guaranteeing optimal paths.

### 3.3 Search vs. MDP Algorithm Comparison

## 4 Conclusion

## 5 References

# Appendices

## A DFS Implementation Code

```
1 class DFSSolver(MazeSolverBase):
2     """
3     Maze solver using Depth-First Search.
4     """
5     def __init__(self, title, maze):
6         """
7         Initialize the maze solver with a given maze array.
8
9         Args:
10         maze (numpy.ndarray): 2D array representing the maze
11         ('#' for walls, '.' for paths, 'S' for start, 'G' for goal).
12         """
13         super().__init__(title, maze)
14
15         self.nodes_explored = 0
16         self.nodes_available = len(list(zip(*np.where(maze == '.'))))
17         self.execution_time = None
18
19         self.start = tuple(zip(*np.where(maze == 'S')))[0]
20         self.goal = tuple(zip(*np.where(maze == 'G')))[0]
21
22         self.visited = np.zeros_like(maze, dtype=bool)
23         self.solution_path = []
24
25     def is_valid_move(self, x, y):
26         """
27         Check if the move is valid (within bounds and not a wall).
28
29         Args:
30         x (int): x-coordinate
31         y (int): y-coordinate
32
33         Returns:
34         bool: True if move is valid, False otherwise.
35         """
36         return (0 <= x < self.width and
37                 0 <= y < self.height and
38                 self.maze[y, x] != '#' and
39                 not self.visited[y, x])
40
41     def solve(self):
42         """
43         Solve the maze using Iterative Depth-First Search.
44
45         Returns:
46         list: Solution path if found, empty list otherwise.
47         """
48         # reset nodes explored, visited array, and solution path
49         self.nodes_explored = 0
50         self.visited = np.zeros_like(self.maze, dtype=bool)
51         self.solution_path = []
52
53         # define start point
54         start_x, start_y = self.start
55
56         # moves for dfs: down, right, up, left
57         moves = [(0, 1), (1, 0), (0, -1), (-1, 0)]
58
59         stack = [(start_x, start_y, [])]
60
61         while stack:
62             current_x, current_y, path = stack.pop()
```



```

63         if self.visited[current_y, current_x]:
64             continue
65
66         self.visited[current_y, current_x] = True
67         self.nodes_explored += 1
68
69         current_path = path + [(current_x, current_y)]
70
71         if (current_x, current_y) == self.goal:
72             self.solution_path = current_path
73             return current_path
74
75         for dx, dy in moves:
76             next_x, next_y = current_x + dx, current_y + dy
77
78             if self.is_valid_move(next_x, next_y):
79                 stack.append((next_x, next_y, current_path))
80
81     # return empty array when no solution found
82     return []
83
84 def get_performance_metrics(self):
85     """
86     Return performance of the DFS algorithm for the maze solution.
87     """
88     return {
89         'path_length': len(self.solution_path),
90         'nodes_explored': self.nodes_explored,
91         'nodes_available': self.nodes_available,
92         'execution_time': self.execution_time,
93         'is_solution_found': bool(self.solution_path)
94     }
95

```

## B BFS Implementation Code

```

1 class BFSSolver(MazeSolverBase):
2     """
3     Maze solver using Breadth-First Search.
4     """
5     def __init__(self, title, maze):
6         """
7         Initialize the maze solver with a given maze array.
8
9         Args:
10         maze (numpy.ndarray): 2D array representing the maze
11         ('#' for walls, '.' for paths, 'S' for start, 'G' for goal).
12         """
13         super().__init__(title, maze)
14
15         self.nodes_explored = 0
16         self.nodes_available = len(list(zip(*np.where(maze == '.'))))
17         self.execution_time = None
18
19         # find start and goal positions
20         self.start = tuple(zip(*np.where(maze == 'S')))[0]
21         self.goal = tuple(zip(*np.where(maze == 'G')))[0]
22
23         self.visited = np.zeros_like(maze, dtype=bool)
24         self.solution_path = []
25
26     def is_valid_move(self, x, y):
27         """
28         Check if the move is valid (within bounds and not a wall).
29
30         Args:
31         x (int): x-coordinate

```

```

32         y (int): y-coordinate
33
34     Returns:
35     bool: True if move is valid, False otherwise.
36     """
37     return (0 <= x < self.width and
38             0 <= y < self.height and
39             self.maze[y, x] != '#' and
40             not self.visited[y, x])
41
42 def solve(self):
43     """
44     Solve the maze using Breadth-First Search.
45
46     Returns:
47     list: Solution path if found, empty list otherwise.
48     """
49     # reset nodes explored, visited array, and solution path
50     self.nodes_explored = 0
51     self.visited = np.zeros_like(self.maze, dtype=bool)
52     self.solution_path = []
53
54     # define start point
55     start_x, start_y = self.start
56
57     # moves for bfs: down, right, up, left
58     moves = [(0, 1), (1, 0), (0, -1), (-1, 0)]
59
60     # use a queue for BFS (First-In-First-Out)
61     queue = deque([(start_x, start_y, [])])
62     self.visited[start_y, start_x] = True # Mark start as visited
63     self.nodes_explored += 1
64
65     while queue:
66         current_x, current_y, path = queue.popleft() # Get the oldest element (FIFO)
67
68         current_path = path + [(current_x, current_y)]
69
70         # Check if goal is reached
71         if (current_x, current_y) == self.goal:
72             self.solution_path = current_path
73             return current_path
74
75         # Try all four directions
76         for dx, dy in moves:
77             next_x, next_y = current_x + dx, current_y + dy
78
79             # If move is valid and cell not visited
80             if self.is_valid_move(next_x, next_y):
81                 queue.append((next_x, next_y, current_path))
82                 self.visited[next_y, next_x] = True # Mark as visited when added to
83                 self.nodes_explored += 1
84
85     # return empty array when no solution found
86     return []
87
88 def get_performance_metrics(self):
89     """
90     Return performance of the BFS algorithm for the maze solution.
91     """
92     return {
93         'path_length': len(self.solution_path),
94         'nodes_explored': self.nodes_explored,
95         'nodes_available': self.nodes_available,
96         'execution_time': self.execution_time,
97         'is_solution_found': bool(self.solution_path)
98     }

```

## C A\* Implementation Code

```
1 class AStarSolver(MazeSolverBase):
2     """
3     Maze solver using A* Search.
4     """
5     def __init__(self, title, maze):
6         """
7         Initialize the maze solver with a given maze array.
8
9         Args:
10            maze (numpy.ndarray): 2D array representing the maze
11            ('#' for walls, '.' for paths, 'S' for start, 'G' for goal).
12            """
13        super().__init__(title, maze)
14
15        self.nodes_explored = 0
16        self.nodes_available = len(list(zip(*np.where(maze == '.'))))
17        self.execution_time = None
18
19        self.counter = 0
20
21        self.start = tuple(zip(*np.where(maze == 'S')))[0]
22        self.goal = tuple(zip(*np.where(maze == 'G')))[0]
23
24        self.visited = np.zeros_like(maze, dtype=bool)
25        self.solution_path = []
26
27    def manhattan_distance(self, pos1, pos2):
28        """
29        Calculate Manhattan distance heuristic.
30
31        Args:
32            pos1 (tuple): First position (x, y)
33            pos2 (tuple): Second position (x, y)
34
35        Returns:
36            int: Manhattan distance between positions
37            """
38        return abs(pos1[0] - pos2[0]) + abs(pos1[1] - pos2[1])
39
40    def is_valid_move(self, x, y):
41        """
42        Check if the move is valid (within bounds and not a wall).
43
44        Args:
45            x (int): x-coordinate
46            y (int): y-coordinate
47
48        Returns:
49            bool: True if move is valid, False otherwise.
50            """
51        return (0 <= x < self.width and
52                0 <= y < self.height and
53                self.maze[y, x] != '#' and
54                not self.visited[y, x])
55
56    def solve(self):
57        """
58        Solve the maze using A* Search.
59
60        Returns:
61            list: Solution path if found, empty list otherwise.
62            """
63        self.nodes_explored = 0
64        self.visited = np.zeros_like(self.maze, dtype=bool)
65        self.solution_path = []
```

```

66
67     # track maximum fringe size for metrics
68     max_fringe_size = 0
69
70     start_pos = self.start
71     goal_pos = self.goal
72
73     parent = {}
74
75     g_score = {start_pos: 0}
76
77     # initialize priority queue for open set
78     # with format: (f_score, counter, position)
79     open_set = [(self.manhattan_distance(start_pos, goal_pos), self.counter, start_pos)]
80     self.counter += 1
81
82     while open_set:
83         # update max fringe size
84         max_fringe_size = max(max_fringe_size, len(open_set))
85
86         # get node with lowest f_score
87         _, _, current = heapq.heappop(open_set)
88         current_x, current_y = current
89
90         # skip if already visited
91         if self.visited[current_y, current_x]:
92             continue
93
94         # mark as visited and update explored nodes
95         self.visited[current_y, current_x] = True
96         self.nodes_explored += 1
97
98         if current == goal_pos:
99             # reconstruct path if goal reached
100             self.solution_path = self._reconstruct_path(parent, current)
101             return self.solution_path
102
103         # explore neighbors in the order: down, right, up, left
104         moves = [(0, 1), (1, 0), (0, -1), (-1, 0)]
105
106         for dx, dy in moves:
107             next_x, next_y = current_x + dx, current_y + dy
108             next_pos = (next_x, next_y)
109
110             if self.is_valid_move(next_x, next_y):
111                 # calculate tentative g-score
112                 tentative_g_score = g_score[current] + 1
113
114                 # when we find a better path to this neighbor
115                 if next_pos not in g_score or tentative_g_score < g_score[next_pos]:
116                     # update the path
117                     parent[next_pos] = current
118                     g_score[next_pos] = tentative_g_score
119
120                     # calculate f_score = g_score + heuristic
121                     f_score = tentative_g_score + self.manhattan_distance(next_pos,
goal_pos)
122
123                     heapq.heappush(open_set, (f_score, self.counter, next_pos))
124                     self.counter += 1
125
126         # return empty array when no path found
127         return []
128
129     def _reconstruct_path(self, parent, current):
130         """
131         Reconstruct the path from start to goal.
132
133         Args:

```

```

133     parent (dict): Dictionary mapping node to its parent
134     current (tuple): Current node (the goal)
135
136     Returns:
137     list: List of positions from start to goal
138     """
139     path = [current]
140     while current in parent:
141         current = parent[current]
142         path.insert(0, current)
143     return path
144
145     def get_performance_metrics(self):
146         """
147         Return performance of the A* algorithm for the maze solution.
148         """
149         return {
150             'path_length': len(self.solution_path),
151             'nodes_explored': self.nodes_explored,
152             'nodes_available': self.nodes_available,
153             'execution_time': self.execution_time,
154             'is_solution_found': bool(self.solution_path)
155         }

```

## D MDP Value Iteration Implementation Code

```

1 class MDPValueIterationSolver(MazeSolverBase):
2     """
3     Maze solver using MDP Value Iteration.
4     """
5     def __init__(self, title, maze, discount_factor=0.9, theta=0.001, max_iterations=1000):
6         """
7         Initialize the MDP Value Iteration solver.
8
9         Args:
10             title: Name/title of the maze
11             maze: The maze to solve (NumPy array)
12             discount_factor: Discount factor for future rewards (gamma)
13             theta: Threshold for determining value convergence
14             max_iterations: Maximum number of iterations to perform
15         """
16         super().__init__(title, maze)
17
18         self.nodes_available = len(list(zip(*np.where(maze == '.'))))
19
20         self.execution_time = None
21
22         self.start = tuple(zip(*np.where(maze == 'S')))[0]
23         self.goal = tuple(zip(*np.where(maze == 'G')))[0]
24
25         self.discount_factor = discount_factor
26         self.theta = theta
27         self.max_iterations = max_iterations
28         self.values = {}
29         self.policy = {}
30         self.iterations = 0
31         self.states_evaluated = 0
32
33         self.solution_path = []
34
35     def is_wall(self, x, y):
36         """Check if a cell is a wall."""
37         # check if coordinates are within bounds
38         if x < 0 or x >= self.width or y < 0 or y >= self.height:
39             return True
40         # check if the cell is a wall (represented by '#')
41         return self.maze[y, x] == '#'

```

```

42
43 def get_states(self):
44     """Get all valid states (positions) in the maze."""
45     states = []
46     for y in range(self.height):
47         for x in range(self.width):
48             # include only non-wall cells as valid states
49             if not self.is_wall(x, y):
50                 states.append((x, y))
51     return states
52
53 def get_actions(self):
54     """Define possible actions as cardinal directions."""
55     return [
56         (-1, 0), # up
57         (1, 0),  # down
58         (0, -1), # left
59         (0, 1)   # right
60     ]
61
62 def get_reward(self, state, next_state):
63     """
64     Define rewards for transitions.
65
66     Args:
67         state: Current state (row, col)
68         next_state: Next state (row, col)
69
70     Returns:
71         reward: Reward for this transition
72     """
73     # set reward = 100 for reaching the goal
74     if next_state == self.goal:
75         return 100
76
77     # set penalty = -100 for hitting a wall
78     if self.is_wall(*next_state):
79         return -100
80
81     # calculate Manhattan distance to goal for getting gradient reward
82     goal_x, goal_y = self.goal
83     next_x, next_y = next_state
84     curr_x, curr_y = state
85
86     curr_dist = abs(curr_x - goal_x) + abs(curr_y - goal_y)
87     next_dist = abs(next_x - goal_x) + abs(next_y - goal_y)
88
89     # if the next state is closer to goal, give bonus
90     if next_dist < curr_dist:
91         return -0.5
92
93     # set small penalty for each step to encourage shorter paths
94     return -1
95
96 def get_transition_prob(self, state, action, next_state):
97     """
98     Get transition probability P(next_state | state, action).
99     For a deterministic environment, this is 1 if next_state is the result of
100     applying action to state, and 0 otherwise.
101     """
102     x, y = state
103     dx, dy = action
104     expected_next_state = (x + dx, y + dy)
105
106     # if next_state is the expected result of the action, probability is 1
107     if expected_next_state == next_state:
108         # check if the move is valid (i.e. it doesn't hit a wall)
109         if not self.is_wall(*expected_next_state):

```

```

110         return 1.0
111
112     # if we're expecting to hit a wall, we stay in the same place
113     if self.is_wall(*expected_next_state) and state == next_state:
114         return 1.0
115
116     # otherwise probability is 0
117     return 0.0
118
119 def solve(self):
120     """
121     Solve the maze using Value Iteration.
122
123     Returns:
124         path: List of positions forming the path from start to goal
125     """
126     print(f"Start position: {self.start}")
127     print(f"Goal position: {self.goal}")
128
129     states = self.get_states()
130     actions = self.get_actions()
131
132     # initialize value function with goal having high value
133     self.values = {state: 0 for state in states}
134     self.values[self.goal] = 100
135
136     # set properties for Value Iteration
137     self.iterations = 0
138     self.states_evaluated = 0
139
140     for i in range(self.max_iterations):
141         self.iterations += 1
142         delta = 0
143
144         # update values for all states
145         for state in states:
146             self.states_evaluated += 1
147
148             # skip updating value if goal state
149             if state == self.goal:
150                 continue
151
152             old_value = self.values[state]
153
154             # calculate new value using Bellman equation
155             new_value = float('-inf')
156
157             for action in actions:
158                 action_value = 0
159                 for next_state in states:
160                     prob = self.get_transition_prob(state, action, next_state)
161
162                     if prob > 0:
163                         reward = self.get_reward(state, next_state)
164                         action_value += prob * \
165                             (reward + self.discount_factor * self.values[next_state])
166                 if action_value > new_value:
167                     new_value = action_value
168
169             self.values[state] = new_value
170             delta = max(delta, abs(old_value - new_value))
171
172         if delta < self.theta:
173             print(f"Value Iteration converged after {i+1} iterations")
174             break
175
176     # extract policy from value function
177     self.policy = {}

```

```

178         for state in states:
179             if state == self.goal:
180                 self.policy[state] = None
181                 continue
182
183             best_action = None
184             best_value = float('-inf')
185
186             for action in actions:
187                 action_value = 0
188
189                 for next_state in states:
190                     prob = self.get_transition_prob(state, action, next_state)
191                     if prob > 0:
192                         reward = self.get_reward(state, next_state)
193                         action_value += prob * \
194                             (reward + self.discount_factor * self.values[next_state])
195
196                 if action_value > best_value:
197                     best_value = action_value
198                     best_action = action
199
200             self.policy[state] = best_action
201
202         path = self.extract_path()
203         self.solution_path = path
204         return path
205
206     def extract_path(self):
207         """
208         Extract the path from start to goal using the computed policy.
209
210         Returns:
211             path: List of positions from start to goal
212         """
213         path = [self.start]
214         current = self.start
215         visited = {self.start} # track visited states to prevent loops
216
217         max_path_length = self.width * self.height
218
219         while current != self.goal and len(path) < max_path_length:
220             action = self.policy[current]
221
222             if action is None:
223                 break
224
225             x, y = current
226             dx, dy = action
227             next_state = (x + dx, y + dy)
228
229             if self.is_wall(*next_state):
230                 break
231
232             # Check for loops
233             if next_state in visited:
234                 print(f"Warning: Loop detected in path at {next_state}")
235                 break
236
237             visited.add(next_state)
238             path.append(next_state)
239             current = next_state
240
241         return path
242
243     def get_performance_metrics(self):
244         """
245         Return performance metrics for the solver.

```



```

246
247     Returns:
248         metrics: Dictionary of performance metrics
249         """
250     path = self.extract_path()
251     is_solution_found = len(path) > 1 and path[-1] == self.goal
252
253     return {
254         'path_length': len(path) if is_solution_found else 0,
255         'nodes_available': self.nodes_available,
256         'iterations': self.iterations,
257         'states_evaluated': self.states_evaluated,
258         'execution_time': self.execution_time,
259         'is_solution_found': is_solution_found
260     }

```

## E MDP Policy Iteration Implementation Code

```

1 class MDPPolicyIterationSolver(MazeSolverBase):
2     """
3     Maze solver using MDP Policy Iteration solver.
4     """
5     def __init__(self, title, maze, discount_factor=0.999,
6                 theta=0.0001, max_iterations=500, policy_eval_iterations=50):
7         """
8         Initialize the MDP Policy Iteration solver.
9
10        Args:
11            title: Name/title of the maze
12            maze: The maze to solve (NumPy array)
13            discount_factor: Discount factor for future rewards (gamma)
14            theta: Threshold for determining value convergence
15            max_iterations: Maximum number of policy iterations to perform
16            policy_eval_iterations: Number of iterations for policy evaluation step
17        """
18        super().__init__(title, maze)
19
20        self.nodes_available = len(list(zip(*np.where(maze == '.'))))
21
22        self.execution_time = None
23
24        self.start = tuple(zip(*np.where(maze == 'S')))[0]
25        self.goal = tuple(zip(*np.where(maze == 'G')))[0]
26
27        self.discount_factor = discount_factor
28        self.theta = theta
29        self.max_iterations = max_iterations
30        self.policy_eval_iterations = policy_eval_iterations
31        self.values = {}
32        self.policy = {}
33        self.iterations = 0
34        self.policy_changes = 0
35        self.states_evaluated = 0
36
37        self.solution_path = []
38
39    def is_wall(self, x, y):
40        """Check if a cell is a wall."""
41        # check if coordinates are within bounds
42        if x < 0 or x >= self.width or y < 0 or y >= self.height:
43            return True
44        # check if the cell is a wall (typically represented by '#')
45        return self.maze[y, x] == '#'
46
47    def get_states(self):
48        """Get all valid states (positions) in the maze."""
49        states = []

```

```

50     for y in range(self.height):
51         for x in range(self.width):
52             # include only non-wall cells as valid states
53             if not self.is_wall(x, y):
54                 states.append((x, y))
55     return states
56
57 def get_actions(self):
58     """Define possible actions as cardinal directions."""
59     return [
60         (-1, 0), # up
61         (1, 0),  # down
62         (0, -1), # left
63         (0, 1)   # right
64     ]
65
66 def get_reward(self, state, next_state):
67     """
68     Enhanced reward structure with goal-directed gradient.
69
70     Args:
71         state: Current state (x, y)
72         next_state: Next state (x, y)
73
74     Returns:
75         reward: Reward for this transition
76     """
77     # set high reward for reaching the goal
78     if next_state == self.goal:
79         return 100
80
81     # set penalty for hitting a wall
82     if self.is_wall(*next_state):
83         return -100
84
85     # calculate Manhattan distance to goal for gradient reward
86     goal_x, goal_y = self.goal
87     next_x, next_y = next_state
88     curr_x, curr_y = state
89
90     # get distances to goal
91     curr_dist = abs(curr_x - goal_x) + abs(curr_y - goal_y)
92     next_dist = abs(next_x - goal_x) + abs(next_y - goal_y)
93
94     # if next state is closer to goal, give bonus
95     if next_dist < curr_dist:
96         return -0.5 # Small step penalty but better than standard step
97
98     # standard step penalty
99     return -1
100
101 def get_transition_prob(self, state, action, next_state):
102     """
103     Get transition probability P(next_state | state, action).
104     For a deterministic environment, this is 1 if next_state is the result of
105     applying action to state, and 0 otherwise.
106     """
107     x, y = state
108     dx, dy = action
109     expected_next_state = (x + dx, y + dy)
110
111     # if next_state is the expected result of the action, probability is 1
112     if expected_next_state == next_state:
113         # check if the move is valid (i.e. it doesn't hit a wall)
114         if not self.is_wall(*expected_next_state):
115             return 1.0
116
117     # if we're expecting to hit a wall, we stay in the same place

```

```

118         if self.is_wall(*expected_next_state) and state == next_state:
119             return 1.0
120
121         # otherwise probability is 0
122         return 0.0
123
124     def policy_evaluation(self, policy, states, actions):
125         """
126         Evaluate policy with convergence check.
127
128         Args:
129             policy: Current policy mapping states to actions
130             states: List of all states
131             actions: List of possible actions
132
133         Returns:
134             values: Dictionary mapping states to their values
135         """
136         values = {state: 0 for state in states}
137
138         # set goal state value higher to create gradient
139         values[self.goal] = 100
140
141         for _ in range(self.policy_eval_iterations):
142             delta = 0
143             for state in states:
144                 self.states_evaluated += 1
145
146                 # skip evaluating if reached goal state
147                 if state == self.goal:
148                     continue
149
150                 old_value = values[state]
151                 action = policy[state]
152
153                 # if no action is defined for this state, skip it
154                 if action is None:
155                     continue
156
157                 new_value = 0
158                 for next_state in states:
159                     prob = self.get_transition_prob(state, action, next_state)
160
161                     if prob > 0:
162                         reward = self.get_reward(state, next_state)
163                         new_value += prob * (reward + self.discount_factor * values[
next_state])
164
165                 values[state] = new_value
166                 delta = max(delta, abs(old_value - new_value))
167
168                 # check for convergence
169                 if delta < self.theta:
170                     break
171
172         return values
173
174     def policy_improvement(self, values, states, actions):
175         """
176         Improve policy based on value function.
177
178         Args:
179             values: Current value function
180             states: List of all states
181             actions: List of possible actions
182
183         Returns:
184             policy: Improved policy

```

```

185         is_stable: Whether the policy has stabilized
186         """
187         policy = {}
188         is_stable = True
189
190         for state in states:
191             if state == self.goal:
192                 policy[state] = None
193                 continue
194
195             old_action = self.policy.get(state)
196
197             best_action = None
198             best_value = float('-inf')
199
200             for action in actions:
201                 action_value = 0
202
203                 for next_state in states:
204                     prob = self.get_transition_prob(state, action, next_state)
205                     if prob > 0:
206                         reward = self.get_reward(state, next_state)
207                         action_value += prob * (reward + self.discount_factor * values[
next_state])
208
209                     if action_value > best_value:
210                         best_value = action_value
211                         best_action = action
212
213                 policy[state] = best_action
214
215                 # check if policy has changed
216                 if old_action != best_action:
217                     is_stable = False
218                     self.policy_changes += 1
219
220             return policy, is_stable
221
222     def solve(self):
223         """
224         Solve the maze using Policy Iteration.
225
226         Returns:
227             path: List of positions forming the path from start to goal
228         """
229         print(f"Start position: {self.start}")
230         print(f"Goal position: {self.goal}")
231
232         states = self.get_states()
233         actions = self.get_actions()
234
235         # initialize policy with goal-directed actions
236         self.policy = {}
237         for state in states:
238             if state == self.goal:
239                 self.policy[state] = None
240             else:
241                 # get all valid actions
242                 valid_actions = []
243                 for action in actions:
244                     x, y = state
245                     dx, dy = action
246                     next_state = (x + dx, y + dy)
247                     if not self.is_wall(*next_state):
248                         valid_actions.append(action)
249
250                 if valid_actions:
251                     # use action that gets closest to goal if possible

```

```

252         goal_x, goal_y = self.goal
253         best_action = None
254         min_distance = float('inf')
255
256         for action in valid_actions:
257             x, y = state
258             dx, dy = action
259             next_x, next_y = x + dx, y + dy
260             dist = abs(next_x - goal_x) + abs(next_y - goal_y)
261
262             if dist < min_distance:
263                 min_distance = dist
264                 best_action = action
265
266         self.policy[state] = best_action
267     else:
268         self.policy[state] = None
269
270     # set properties for policy iteration
271     self.iterations = 0
272     self.policy_changes = 0
273     self.states_evaluated = 0
274
275     for i in range(self.max_iterations):
276         self.iterations += 1
277
278         # one, do Policy Evaluation
279         self.values = self.policy_evaluation(self.policy, states, actions)
280
281         # two, do Policy Improvement
282         new_policy, is_stable = self.policy_improvement(self.values, states, actions)
283         self.policy = new_policy
284
285         # now check if policy has stabilized
286         if is_stable:
287             print(f"Policy Iteration converged after {i+1} iterations")
288             break
289
290     path = self.extract_path()
291     self.solution_path = path
292     return path
293
294     def extract_path(self):
295         """
296         Extract the path from start to goal using the computed policy.
297
298         Returns:
299             path: List of positions from start to goal
300         """
301         path = [self.start]
302         current = self.start
303         visited = {self.start} # track visited states to prevent loops
304
305         max_path_length = self.width * self.height
306
307         while current != self.goal and len(path) < max_path_length:
308             action = self.policy[current]
309
310             if action is None:
311                 break
312
313             x, y = current
314             dx, dy = action
315             next_state = (x + dx, y + dy)
316
317             # check if the next state is valid
318             if self.is_wall(*next_state):
319                 # this shouldn't happen with a valid policy

```

```

320         break
321
322     # Check for loops
323     if next_state in visited:
324         print(f"Warning: Loop detected in path at {next_state}")
325         break
326
327     visited.add(next_state)
328     path.append(next_state)
329     current = next_state
330
331     return path
332
333 def get_performance_metrics(self):
334     """
335     Return performance metrics for the solver.
336
337     Returns:
338         metrics: Dictionary of performance metrics
339     """
340     path = self.extract_path()
341     is_solution_found = len(path) > 1 and path[-1] == self.goal
342
343     return {
344         'path_length': len(path) if is_solution_found else 0,
345         'nodes_available': self.nodes_available,
346         'iterations': self.iterations,
347         'policy_changes': self.policy_changes,
348         'states_evaluated': self.states_evaluated,
349         'execution_time': self.execution_time,
350         'is_solution_found': is_solution_found
351     }

```